

Bokun Wang

bokun.wang@utexas.edu | +1 319-512-3088

Research Interests

Resource-Efficient Generative AI, Edge AI, AI Systems

Research and Professional Experience

Postdoctoral Researcher, The University of Texas at Austin

September 2025–present

Chandra Family Department of Electrical and Computer Engineering

Faculty host: Diana Marculescu

Machine Learning Intern, Arm

May–August 2023

Central Technology / Machine Learning Group, Boston, MA

Education

Texas A&M University

2021–2025

Ph.D. in Computer Science, Advisor: Tianbao Yang

College Station, Texas

University of California, Davis

2018–2020

M.S. in Computer Science

Davis, California

University of Electronic Science and Technology of China

2014–2018

B.S. in Computer Science, Yingcai Honors College Program

Chengdu, China

Honors and Awards

2026 Google Research Award (Awarded to PI Diana Marculescu; I was the primary proposal writer)

For Machine Learning Research and Education with TPUs (Acceptance rate: 10%)

2026 Early Career Postdoctoral Fellow, The University of Texas at Austin

2026 FutureBAProf, Tippie College of Business, The University of Iowa

22 selected from 117 applicants

2026 Gold Reviewer Award, ICML

2025 Notable Reviewer, ICLR

2024 Oral Presentation, Conference on the Mathematical Theory of Deep Neural Networks (DeepMath)

2022 Scholar Award, NeurIPS

2022 Travel Award, ICML

2022 Spotlight Presentation, ICLR

2018 Provincial Outstanding Graduate, Sichuan Province, China

2016 National Scholarship, Ministry of Education of the People's Republic of China

2015 People's Scholarship (Highest Class), University of Electronic Science and Technology of China

2014 Freshman Scholarship, University of Electronic Science and Technology of China

Teaching Experience

Guest Instructor & Course Materials Development, Texas A&M University

Fall 2023

CSCE 689: Optimization for Machine Learning (Now CSCE 726)

I delivered guest lectures and helped prepare lecture notes that were later expanded into the forthcoming Springer book, *Compositional Optimization for Advanced Machine Learning*.

Teaching Assistant, University of California, Davis

Winter 2020

ECS 32B: Introduction to Data Structures

Teaching Assistant, University of California, Davis

Fall 2019

ECS 154A: Computer Architecture

Teaching Assistant, University of California, Davis

Spring 2019

ECS 170: Introduction to Artificial Intelligence

Teaching Assistant & Course Materials Development, University of California, Davis

Winter 2019

ECS 271: Machine Learning and Data Discovery I redesigned all course projects and in-class software demos to reflect the latest advances in the field, which earned positive feedback from students.

Manuscripts and Technical Reports

1. **Bokun Wang**, et al. "ACSeL: Adaptive Configuration Search for Hardware-Constrained LLM Inference." Under review, 2026.
2. Natalia Frumkin, **Bokun Wang**, et al. "DARE: Diffusion Language Model Activation Reuse for Efficient Inference." Under review, 2026.
3. Hung-Yueh Chiang, **Bokun Wang**, Diana Marculescu. "ELANA: A Simple Energy and Latency Analyzer for LLMs." Technical report, 2026.

Selected Recent Publications (Full List)

1. Dennis Menn, Yuedong Yang, **Bokun Wang**, et al. "Training-free Latent Inter-Frame Pruning with Attention Recovery." *Conference on Neural Information Processing Systems (NeurIPS)*, 2026.
2. Xiyuan Wei, Linli Zhou, **Bokun Wang**, et al. "A Geometry-Aware Efficient Algorithm for Compositional Entropic Risk Minimization." *International Conference on Machine Learning (ICML)*, 2026.
3. Vicente Balmaseda, **Bokun Wang**, et al. "Discovering Global False Negatives On the Fly for Self-supervised Contrastive Learning." *International Conference on Machine Learning (ICML)*, 2025.
4. **Bokun Wang**, et al. "A Near-Optimal Single-Loop Algorithm for Convex Finite-Sum Coupled Compositional Stochastic Optimization." *International Conference on Machine Learning (ICML)*, 2025.
5. **Bokun Wang**, et al. "On Discriminative Probabilistic Modeling for Self-Supervised Representation Learning." *International Conference on Learning Representations (ICLR)*, 2025.
6. **Bokun Wang**, et al. "Memory-Based Optimization Methods for Model-Agnostic Meta-Learning and Personalized Federated Learning." *Journal of Machine Learning Research (JMLR)*, 2023.
7. Konstantin Mishchenko, **Bokun Wang**, et al. "IntSGD: Adaptive Floatless Compression of Stochastic Gradients." *International Conference on Learning Representations (ICLR)*, 2022.
8. **Bokun Wang**, et al. "Finite-Sum Coupled Compositional Stochastic Optimization Theory and Applications." *International Conference on Machine Learning (ICML)*, 2022.
9. **Bokun Wang**, et al. "Riemannian Stochastic Proximal Gradient Methods for Nonsmooth Optimization over the Stiefel Manifold." *Journal of Machine Learning Research (JMLR)*, 2022.

Service

Organizer of session "New Min-max Optimization Algorithms and Their Applications in Machine Learning and AI", 2024 INFORMS Annual Meeting, Seattle

Reviewer of *Mathematical Programming*, *SIAM Journal on Optimization*, *SIAM Review*, *Journal of Machine Learning Research*, *Transactions of Machine Learning Research*, *IJSE Transactions*, *ICML*, *NeurIPS*, *ICLR*, *AISTATS*

References

1. **Diana Marculescu**
Department Chair and Motorola Regents Chair Professor
Chandra Family Department of Electrical and Computer Engineering (ECE), University of Texas at Austin
dianam@utexas.edu
2. **Tianbao Yang**
Professor and Stephen Horn '79 Engineering Excellence Chair
Department of Computer Science and Engineering, Texas A&M University
tianbao-yang@tamu.edu
3. **Radu Marculescu**
Laura Jennings Turner Chair in Engineering
Chandra Family Department of Electrical and Computer Engineering (ECE), University of Texas at Austin
radum@utexas.edu
4. **Qihang Lin**
Gary C. Fethke Research Professor in Business Analytics
Department of Business Analytics, University of Iowa

qihang-lin@uiowa.edu