

Bokun Wang

bokun.wang@utexas.edu | +1 319-512-3088

Research Interests

Resource-Efficient Generative AI, Machine Learning.

Education

Texas A&M University	2021–2025
Ph.D. in Computer Science	College Station, Texas
Advisor: Tianbao Yang	
University of California, Davis	2018–2020
M.S. in Computer Science	Davis, California
University of Electronic Science and Technology of China	2014–2018
B.S. in Computer Science, Yingcai Honors College Program	Chengdu, China

Research and Professional Experience

Postdoctoral Researcher, The University of Texas at Austin	September 2025–present
Chandra Family Department of Electrical and Computer Engineering	
Faculty host: Diana Marculescu	
Machine Learning Intern, Arm	May–August 2023
Central Technology / Machine Learning Group	

Honors and Awards

2026	Early Career Postdoctoral Fellow, The University of Texas at Austin
2026	FutureBAProf, Tippie College of Business, The University of Iowa 22 selected from 117 applicants
2026	Gold Reviewer Award, ICML
2025	Notable Reviewer, ICLR
2024	Oral Presentation, Conference on the Mathematical Theory of Deep Neural Networks (DeepMath)
2022	Scholar Award, NeurIPS
2022	Travel Award, ICML
2022	Spotlight Presentation, ICLR
2018	Provincial Outstanding Graduate, Sichuan Province, China
2017	Lixin Tang Scholarship, University of Electronic Science and Technology of China
2016	National Scholarship, Ministry of Education of the People's Republic of China
2015	People's Scholarship (Special Class), University of Electronic Science and Technology of China
2014	Freshman Scholarship, University of Electronic Science and Technology of China

Teaching Experience

Guest Instructor, Texas A&M University	Fall 2023
CSCE 689: Optimization for Machine Learning	
Teaching Assistant, University of California, Davis	Winter 2020
ECS 32B: Introduction to Data Structures	
Teaching Assistant, University of California, Davis	Fall 2019
ECS 154A: Computer Architecture	
Teaching Assistant, University of California, Davis	Spring 2019
ECS 170: Introduction to Artificial Intelligence	
Teaching Assistant, University of California, Davis	Winter 2019
ECS 271: Machine Learning and Data Discovery	

Manuscripts and Technical Reports

1. Bokun Wang, Hung-Yueh Chiang, Diana Marculescu. "ACSeL: Adaptive Configuration Search for Hardware-

- Constrained LLM Inference.” Under review, 2026.
2. Natalia Frumkin, **Bokun Wang**, et al. “DARE: Diffusion Language Model Activation Reuse for Efficient Inference.” Under review, 2026.
 3. Hung-Yueh Chiang, **Bokun Wang**, Diana Marculescu. “ELANA: A Simple Energy and Latency Analyzer for LLMs.” Technical report, 2026.

Selected Recent Publications (Full List)

1. Dennis Menn, Yuedong Yang, **Bokun Wang**, et al. “Training-free Latent Inter-Frame Pruning with Attention Recovery.” Conference on Neural Information Processing Systems (NeurIPS), 2026.
2. Xiyuan Wei, Linli Zhou, **Bokun Wang**, et al. “A Geometry-Aware Efficient Algorithm for Compositional Entropic Risk Minimization.” *International Conference on Machine Learning (ICML)*, 2026.
3. Vicente Balmaseda, **Bokun Wang**, et al. “Discovering Global False Negatives On the Fly for Self-supervised Contrastive Learning.” *International Conference on Machine Learning (ICML)*, 2025.
4. **Bokun Wang**, et al. “A Near-Optimal Single-Loop Algorithm for Convex Finite-Sum Coupled Compositional Stochastic Optimization.” *International Conference on Machine Learning (ICML)*, 2025.
5. **Bokun Wang**, et al. “On Discriminative Probabilistic Modeling for Self-Supervised Representation Learning.” *International Conference on Learning Representations (ICLR)*, 2025.
6. **Bokun Wang**, et al. “Memory-Based Optimization Methods for Model-Agnostic Meta-Learning and Personalized Federated Learning.” *Journal of Machine Learning Research (JMLR)*, 2023.
7. Konstantin Mishchenko, **Bokun Wang**, et al. “IntSGD: Adaptive Floatless Compression of Stochastic Gradients.” *International Conference on Learning Representations (ICLR)*, 2022.
8. **Bokun Wang**, et al. “Finite-Sum Coupled Compositional Stochastic Optimization Theory and Applications.” *International Conference on Machine Learning (ICML)*, 2022.
9. **Bokun Wang**, et al. “Riemannian Stochastic Proximal Gradient Methods for Nonsmooth Optimization over the Stiefel Manifold.” *Journal of Machine Learning Research (JMLR)*, 2022.