

# Bokun Wang

bokun.wang@utexas.edu | +1 319-512-3088

## Research Interests

---

Resource-Efficient Generative AI, Edge AI, AI Systems

## Research and Professional Experience

---

**Postdoctoral Researcher, The University of Texas at Austin**

September 2025–present

Chandra Family Department of Electrical and Computer Engineering  
Faculty host: Diana Marculescu

**Machine Learning Intern, Arm**

May–August 2023

Central Technology / Machine Learning Group, Boston, MA

## Education

---

**Texas A&M University**

2021–2025

Ph.D. in Computer Science, Advisor: Tianbao Yang

College Station, Texas

**University of California, Davis**

2018–2020

M.S. in Computer Science

Davis, California

**University of Electronic Science and Technology of China**

2014–2018

B.S. in Computer Science, Yingcai Honors College Program

Chengdu, China

## Honors and Awards

---

**2026** Google Research Award (Awarded to PI Diana Marculescu; I served as the primary proposal writer)

**2026** Early Career Postdoctoral Fellow, The University of Texas at Austin

**2026** FutureBAProf, Tippie College of Business, The University of Iowa  
22 selected from 117 applicants

**2026** Gold Reviewer Award, ICML

**2025** Notable Reviewer, ICLR

**2024** Oral Presentation, Conference on the Mathematical Theory of Deep Neural Networks (DeepMath)

**2022** Scholar Award, NeurIPS

**2022** Travel Award, ICML

**2022** Spotlight Presentation, ICLR

**2018** Provincial Outstanding Graduate, Sichuan Province, China

**2016** National Scholarship, Ministry of Education of the People's Republic of China

**2015** People's Scholarship (Highest Class), University of Electronic Science and Technology of China

**2014** Freshman Scholarship, University of Electronic Science and Technology of China

## Teaching Experience

---

**Guest Instructor & Course Materials Development, Texas A&M University**

Fall 2023

CSCE 689: Optimization for Machine Learning (Now CSCE 726)

I delivered guest lectures and helped prepare lecture notes that were later expanded into the forthcoming Springer book, *Compositional Optimization for Advanced Machine Learning*.

**Teaching Assistant, University of California, Davis**

Winter 2020

ECS 32B: Introduction to Data Structures

**Teaching Assistant, University of California, Davis**

Fall 2019

ECS 154A: Computer Architecture

**Teaching Assistant, University of California, Davis**

Spring 2019

ECS 170: Introduction to Artificial Intelligence

**Teaching Assistant & Course Materials Development, University of California, Davis**

Winter 2019

ECS 271: Machine Learning and Data Discovery I redesigned all course projects and in-class software demos to reflect the latest advances in the field, which earned positive feedback from students.

## Manuscripts and Technical Reports

- 
1. **Bokun Wang**, et al. "ACSeL: Adaptive Configuration Search for Hardware-Constrained LLM Inference." Under review, 2026.
  2. Natalia Frumkin, **Bokun Wang**, et al. "DARE: Diffusion Language Model Activation Reuse for Efficient Inference." Under review, 2026.
  3. Hung-Yueh Chiang, **Bokun Wang**, Diana Marculescu. "ELANA: A Simple Energy and Latency Analyzer for LLMs." Technical report, 2026.

## Selected Recent Publications (Full List)

---

1. Dennis Menn, Yuedong Yang, **Bokun Wang**, et al. "Training-free Latent Inter-Frame Pruning with Attention Recovery." *Conference on Neural Information Processing Systems (NeurIPS)*, 2026.
2. Xiyuan Wei, Linli Zhou, **Bokun Wang**, et al. "A Geometry-Aware Efficient Algorithm for Compositional Entropic Risk Minimization." *International Conference on Machine Learning (ICML)*, 2026.
3. Vicente Balmaseda, **Bokun Wang**, et al. "Discovering Global False Negatives On the Fly for Self-supervised Contrastive Learning." *International Conference on Machine Learning (ICML)*, 2025.
4. **Bokun Wang**, et al. "A Near-Optimal Single-Loop Algorithm for Convex Finite-Sum Coupled Compositional Stochastic Optimization." *International Conference on Machine Learning (ICML)*, 2025.
5. **Bokun Wang**, et al. "On Discriminative Probabilistic Modeling for Self-Supervised Representation Learning." *International Conference on Learning Representations (ICLR)*, 2025.
6. **Bokun Wang**, et al. "Memory-Based Optimization Methods for Model-Agnostic Meta-Learning and Personalized Federated Learning." *Journal of Machine Learning Research (JMLR)*, 2023.
7. Konstantin Mishchenko, **Bokun Wang**, et al. "IntSGD: Adaptive Floatless Compression of Stochastic Gradients." *International Conference on Learning Representations (ICLR)*, 2022.
8. **Bokun Wang**, et al. "Finite-Sum Coupled Compositional Stochastic Optimization Theory and Applications." *International Conference on Machine Learning (ICML)*, 2022.
9. **Bokun Wang**, et al. "Riemannian Stochastic Proximal Gradient Methods for Nonsmooth Optimization over the Stiefel Manifold." *Journal of Machine Learning Research (JMLR)*, 2022.

## Service

---

Organizer of session "New Min-max Optimization Algorithms and Their Applications in Machine Learning and AI", 2024 INFORMS Annual Meeting, Seattle

Reviewer of *Mathematical Programming*, *SIAM Journal on Optimization*, *SIAM Review*, *Journal of Machine Learning Research*, *Transactions of Machine Learning Research*, *IJSE Transactions*, *ICML*, *NeurIPS*, *ICLR*, *AISTATS*

## References

---

1. **Diana Marculescu**  
Department Chair and Motorola Regents Chair Professor  
Chandra Family Department of Electrical and Computer Engineering (ECE), University of Texas at Austin  
dianam@utexas.edu
2. **Tianbao Yang**  
Professor and Stephen Horn '79 Engineering Excellence Chair  
Department of Computer Science and Engineering, Texas A&M University  
tianbao-yang@tamu.edu
3. **Radu Marculescu**  
Laura Jennings Turner Chair in Engineering  
Chandra Family Department of Electrical and Computer Engineering (ECE), University of Texas at Austin  
radum@utexas.edu
4. **Qihang Lin**  
Gary C. Fethke Research Professor in Business Analytics  
Department of Business Analytics, University of Iowa  
qihang-lin@uiowa.edu